

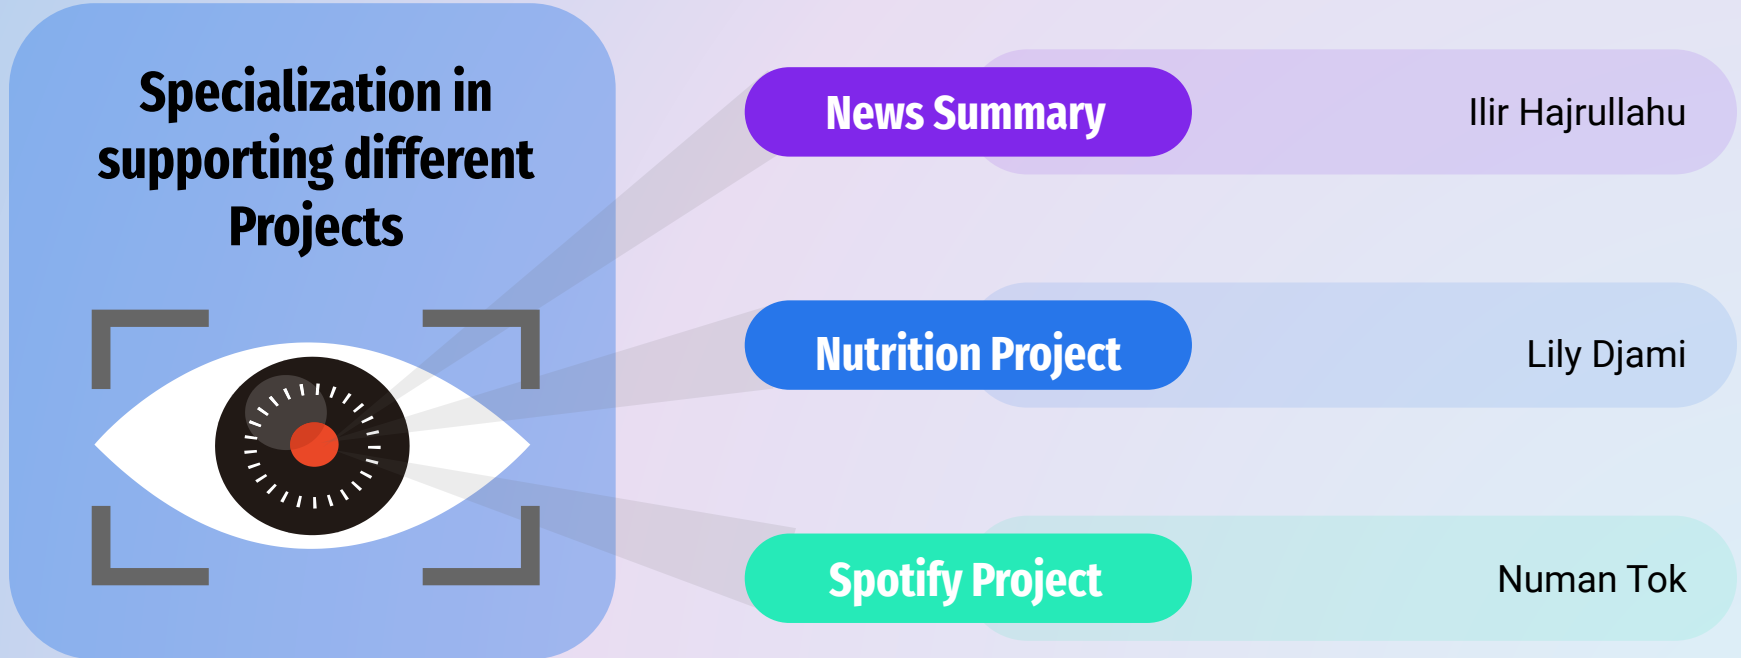
# Machine Learning Team


Lily Djami, Numan Tok

# Overview



# Team Structure and Responsibilities





# Nutrition Group Object Detection

Machine Learning Model

# Requirement Questions for Nutrition Project

## General Questions:

- What dataset are you going to use?
- What is the expected input?
- How accurate should the model be?
- What is the output that you want, what do you want to predict?
- What is the desired speed of the output creation?
- How do you judge if a predicted output is good or bad?

## Specific for Nutrition Project:

- Is there enough image data, or would simulated data be needed?
- Would the model be used to only detect food from a specific cuisine (Asian/Western/etc.)?
- Will there be multiple objects (different fruits) in the picture or just one? Do you want to detect them one after another or as a group?

# Requirements - Nutrition Group

- **Problem:** detection of food ingredients from an image (object detection)
- **Input:** Image with multiple ingredients
- **Output:** List of detected ingredients
- **Good prediction:** all food items are correctly identified.
- **Accuracy:** 90%, but user can delete bad predictions
- **Speed:** less than 1 minute
- **Type of food:** no information
- **Datasets:** food image datasets based on initial results (fruits and vegetables, sausages, bread, grocery items)



Figure: Sample input image

# System Architecture - Faster R-CNN

- One of the most popular class of algorithms for object detection
- Works similarly to CNN in image classification, however the classification is region based.
- Searches through image for possible regions of interests.
- Extract regions using their bounding boxes and classify using CNN.
- Faster R-CNN the latest version with faster region search algorithm

# System Architecture - Faster RCNN

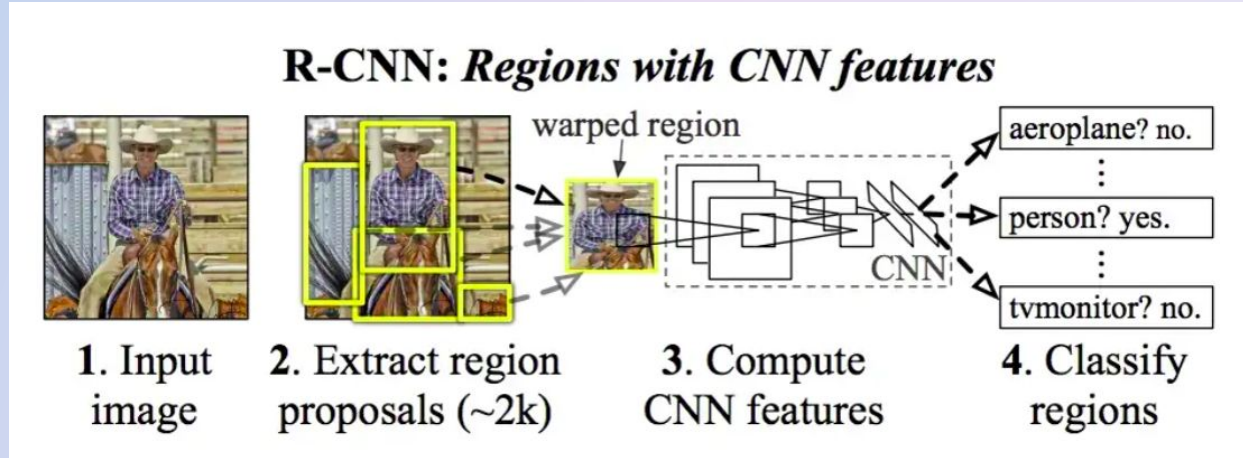


Figure: The R-CNN architecture



# Initial Results



- Base model: Faster R-CNN trained with ResNet-50-FPN from PyTorch, minimum size 4096
  - Detected: cheese, orange (lemon), egg
  - Not detected: milk, bread, peppers
- Retraining to improve results

Figure: Initial results. Items on the top row was detected, but not items on the bottom row

# Datasets for Retraining

After discussion with the nutrition group, the following datasets are chosen:

- Sausage
- Bread
- Fruits and vegetables
- Freiburg Groceries

Training:

- Training on Google Colab
- Label set collected from all datasets (66 labels)
- The same label set used for all training
- Base model trained on 50 epochs/dataset



Figure: Sample images from the datasets.

# Results

Retraining resulted in worse results

- Mislabeling, less detected objects

Possible causes:

- Label set misalignment
- Could be fixed with further re-training with correct label set, but there was no time.
- Used the rest of the time to help integration with Nutrition group.





Figure: Sample images from the datasets.



# Spotify Group Recommendation

Machine Learning Model

# Requirement Questions for Spotify Project

## General Questions:

- What dataset are you going to use?
- What is the expected input?
- How accurate should the model be?
- What is the output that you want, what do you want to predict?
- What is the desired speed of the output creation?
- How do you judge if a predicted output is good or bad?

## Specific for Nutrition Project:

- Do you want a feedback loop to the model?
- Which features do you want to use for the prediction? Length of song, mood, ...?
- Do we have access to the Blend feature functions to possibly build on? (e.g. in the Spotify API)
- Do we also have access to functions of the "Recommended" feature?
- Is there a group size limit or other restrictions?

# Requirements - Spotify Group

- **Problem:** Create a group playlist that satisfies the music taste of all group members.
- **Input:** Can be defined by ML Group. Spotify can format/parse it accordingly etc.
- **Output:** List of matching tracks (Track IDs).
- **Good prediction:** No specification
- **Accuracy:** No specification
- **Speed:** No specification
- **Feedback loop:** Users will be able to rate playlists
- **Features:** Free to choose for ML group.
- **Blend feature:** Does not seem to be available.
- **"Recommended" feature:** Available
- **Limits or other restrictions:** No specification
- **Datasets:** Data about the songs and the playlists (from API). Among other things, genre, abstract popularity info and some user-related information are available.

# Requirement Questions - Spotify Group

- What dataset are you going to use?
  - Via the Spotify API we get some data about the songs and the playlists.  
Among other things, genre, abstract popularity info and some user-related information.
- What is the expected input?
  - Can be defined by ML Group. Spotify can format/parse it accordingly etc.
- How do you judge if a predicted output is good or bad?
  - No specifications.
- How accurate should the model be? Are there outputs that would be discarded at a specific limit (no result for the user)?
  - No specifications.
- What is the output that you want, what do you want to predict?
  - Provide a list of matching songs. The Spotify Group will then create a playlist via the Spotify API.
- What is the desired speed of the output creation?
  - No specifications.

# Requirement Questions - Spotify Group

- How do you judge the created playlist if it is accurate or not?  
→ No idea for testing phase. Can be tested in production by user feedback or playlist listening counts etc.
- Do you want a feedback loop to the model?  
→ There should be feedback from the group. A direct rating from the users was planned.
- Which features do you want to use for the prediction? Length of song, mood, ...?  
→ Free to choose for ML group.
- Do we have access to the Blend feature functions to possibly build on? (e.g. in the Spotify API)  
→ Does not seem to be available.
- Do we also have access to functions of the "Recommended - Based on the content in this playlist" feature in order to use it in our code?  
→ Yes, there is an API Endpoint for the recommendation feature.
- Is there a group size limit or other restrictions?  
→ No.



# Understanding the Fundamentals

Problem:

Create a group playlist that satisfies the music taste of all group members.

To solve this, we should think about:

- What is a Music Track?
- What is a Playlist?

# What is a Music Track?

- 1) Instruments: A track usually includes different kinds of instruments or electronic sounds that form the basis of a track.
- 2) Vocals: A track can contain vocals from one or many persons that may be detailed lyrics or just single words.
- 3) Effects: Auto tune, echos, fade in/out, delays,...
- 4) Tempo: The track speed that is measured in BPM (beats per minute).
- 5) Composition: A track may include several parts with different characteristics
- 6) Mode, Key & time signatures: Track features that influence the overall mood
- 7) ... Has many characteristics

# What is a Music Playlist?

- 1) A sequence of tracks
- 2) In most cases including tracks with similar characteristics
- 3) The tracks may or may not be in a harmonic order  
(Sorted in a way that 2 consecutive tracks have a low difference of their characteristics)

# Data Fundament & API possibilities

- Get Track's Audio Features

- "danceability": 0.696
- "energy": 0.905
- "key": 2
- "loudness": -2.743
- "mode": 1
- "acousticness": 0.011
- "instrumentalness": 0.000905
- "liveness": 0.302
- "valence": 0.625
- "tempo": 114.944
- "duration\_ms": 207960
- "time\_signature": 4

# Data Fundament & API possibilities

- Get Track's Audio Features

- "danceability": 0.696
- "energy": 0.905
- "key": 2
- "loudness": -2.743
- "mode": 1
- "acousticness": 0.011
- "instrumentalness": 0.000905
- "liveness": 0.302
- "valence": 0.625
- "tempo": 114.944
- "duration\_ms": 207960
- "time\_signature": 4

- Get Track

- Release\_date
- Duration
- Popularity

- Get Track's Audio Analysis

- Offset\_seconds
- Window\_seconds
- End\_of\_fade\_in / Start\_of\_fade\_out
- ... More detailed info about some audio features

- Get user's saved tracks: (limit 20)

- Get a list of Spotify featured playlists (limit 50)

- Get user's playlists: (limit 50)

- Get user's top tracks : (limit 20)

# Solution 1 – Clustering + ANN (hybrid model)

## Clustering

- Try to find the most common similarities of the tracks that the different group members listen to:
  1. Group tracks into clusters based on the similarity of the audio features
  2. Find the most promising cluster to use it as an “ideal characteristics representative” by statistical criteria like
    - Percentage of covered tracks
    - Balancing of user representation
    - Additional ranking of tracks (e.g. by recency of when the track was saved to the library; initially a use of listening counts was planned)


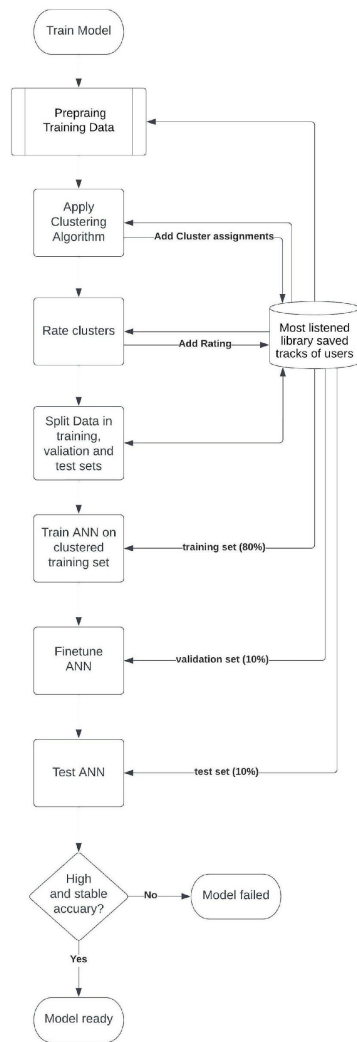
# Solution 1 – Clustering + ANN (hybrid model)

## Clustering

- Try to find the most common similarities of the tracks that the different group members listen to:
  1. Group tracks into clusters based on the similarity of the audio features
  2. Find the most promising cluster to use it as an “ideal characteristics representative” by statistical criteria like
    - Percentage of covered tracks
    - Balancing of user representation
    - Additional ranking of tracks (e.g. by recency of when the track was saved to the library; initially a use of listening counts was planned)

## ANN

- Train an ANN to predict if a song matches the “ideal characteristics representative”
  1. Use the track’s audio features and cluster labels as training input
  2. Predict the class of new songs
- The higher the predicted value for the most promising cluster, the more suitable it is for the playlist





## Solution 2 – Graph model

□ Find out how the audio features of consecutive tracks change in the playlists of all group members. Determine a weight vector as importance measure for the audio features:

1. Calculate the differences between the audio features of each pair of consecutive tracks
2. Sum the differences per audio feature = Vector of summed feature distances (“overall\_distances”)
3. A lower difference is interpreted as a higher importance  $\Rightarrow$  get a higher weight  
(feature\_weights =  $1 / \text{overall\_distances}$ )

## Solution 2 – Graph model

- Find out how the audio features of consecutive tracks change in the playlists of all group members. Determine a weight vector as importance measure for the audio features:
  1. Calculate the differences between the audio features of each pair of consecutive tracks
  2. Sum the differences per audio feature = Vector of summed feature distances (“overall\_distances”)
  3. A lower difference is interpreted as a higher importance  $\Rightarrow$  get a higher weight  
(feature\_weights =  $1 / \text{overall\_distances}$ )
  
- Create an undirected complete graph for all selectable tracks with
  - Nodes = Tracks
  - Edge weights = Weighted difference between the audio features of the connected Nodes

# Solution 2 – Graph model

- Find out how the audio features of consecutive tracks change in the playlists of all group members. Determine a weight vector as importance measure for the audio features:
  1. Calculate the differences between the audio features of each pair of consecutive tracks
  2. Sum the differences per audio feature = Vector of summed feature distances (“overall\_distances”)
  3. A lower difference is interpreted as a higher importance  $\Rightarrow$  get a higher weight  
(feature\_weights =  $1 / \text{overall\_distances}$ )
  
- Create an undirected complete graph for all selectable tracks with
  - Nodes = Tracks
  - Edge weights = Weighted difference between the audio features of the connected Nodes
  
- Create a playlist by finding paths through the graph
  1. Specify a set of starting points: For each user, select his/her 3 top tracks (available through Spotify API)
  2. For each start track, find a sub-playlist (aka path) by successively adding the nearest neighbor in the graph
  3. Concatenate the sub-playlists that together form the overall result

# Model Comparison

## Clustering + ANN

- Higher complexity / more parts
- More influencing parameters that need to be carefully chosen (e.g. selection of Clustering algorithm, Sizes and numbers of layers)
- Depends on well-chosen cluster goodness measure
- Depends on finding a good cluster
- Lower Transparency of Decisioning
- Higher computational cost

# Model Comparison

## Clustering + ANN

- Higher complexity / more parts
- More influencing parameters that need to be carefully chosen (e.g. selection of Clustering algorithm, Sizes and numbers of layers)
- Depends on well-chosen cluster goodness measure
- Depends on finding a good cluster
- Lower Transparency of Decisioning
- Higher computational cost

## Graph model

- Quality of playlist depends on start songs (but 3 top songs for each should be a good basis)
- There may be interruptions of the playlist flow because of the sub-playlist approach
- There could be too much similarity of the tracks
  - Could be countered by adding random noise, prefiltering the track universe or other techniques

# Testing

- Unit Tests for each method
- Integration Tests: for combinations of methods, modules and classes and for the whole ML part
- Static testing: walkthroughs with the Spotify project team
- Performance Tests
  - optimization through:
    1. Replacing for loops by list comprehensions
    2. Minimizing API calls by retrieving the audio features and playlists for whole chunks at once (API limits of 100 tracks and 50 playlists)

# Runtime Optimization

```
--- retrieve_audio_features: 1.6789839267730713 seconds ---  
--- retrieve_audio_features: 2.0816640853881836 seconds ---  
--- retrieve_audio_features: 1.311901569366455 seconds ---  
--- retrieve_audio_features: 3.132132053375244 seconds ---  
--- retrieve_audio_features: 4.034107208251953 seconds ---  
--- retrieve_audio_features: 2.705784797668457 seconds ---  
--- retrieve_audio_features: 3.4747555255889893 seconds ---  
--- custom_audio_features_playlists: 142.38457870483398 seconds ---  
--- prepare_user_playlists: 142.10485696792603 seconds ---
```

```
--- retrieve_audio_features: 0.08209395408630371 seconds ---  
--- retrieve_audio_features: 0.12010526657104492 seconds ---  
--- retrieve_audio_features: 0.09958815574645996 seconds ---  
--- retrieve_audio_features: 0.0886235237121582 seconds ---  
--- retrieve_audio_features: 0.10134363174438477 seconds ---  
--- retrieve_audio_features: 0.2505049705505371 seconds ---  
--- retrieve_audio_features: 0.1434917449951172 seconds ---  
--- custom_audio_features_playlists: 13.96532940864563 seconds ---  
--- prepare_user_playlists: 14.185040473937988 seconds ---
```

## 1. Optimization

- On retrieve\_audio\_features
- -90% runtime

## 2. Optimization

- On prepare\_track\_universe
- -64% runtime

```
ode\ML model\graph-solution> python ml_main.py  
--- retrieve_audio_features: 0.12012146949768066 seconds ---  
--- prepare_track_universe: 19.110339822769165 seconds ---
```

```
model\graph-solution> python ml_main.py  
--- retrieve_audio_features: 0.12018020915985107 seconds ---  
--- prepare_track_universe: 6.9215781688690186 seconds ---
```

# TODOS Numan

- Create Flow Chart for solution 2 and insert
- Split slides & optimize slides
  - Fix Duplicates in Requirements